

# Bayesian Learning - Part A

Junqi Hu

2025-09-30

## Table of contents

Problem 1 . . . . .	1
Problem 1a) . . . . .	1
Problem 1b) . . . . .	4
Problem 1c) . . . . .	4
Problem 2 . . . . .	5
Problem 2a) . . . . .	6
Problem 2b) . . . . .	7
Problem 2c) . . . . .	8
Problem 2d) . . . . .	9
Problem 2e) . . . . .	10
Problem 3 . . . . .	11
Problem 3a) . . . . .	11
Problem 3b) . . . . .	12
Problem 3c) . . . . .	14
Problem 4 . . . . .	15
Problem 4a) Determine a suitable prior distribution . . . . .	16
Problem 4b) Simulating from the posterior . . . . .	20
Problem 4c) Locating the day with the highest expected temperature . . . . .	24

## Problem 1

### Problem 1a)

In this task, we assume a  $Beta(\alpha_0, \beta_0)$  prior for  $\theta$  which comes from  $y_1, \dots, y_n | \theta \sim Bern(\theta)$ . We use Monte Carlo methods to estimate the posterior and standard deviation.

```

set.seed(42)
n = 20
s = 14
f = n-s
alpha0 = 2
beta0 = 2

# Use posterior formula
alpha_post = alpha0+s
beta_post = beta0+f

# Set true value as benchmark
mean_true = alpha_post/(alpha_post+beta_post)
var_true = (alpha_post*beta_post)/(((alpha_post+beta_post)**2)
                                     * (alpha_post+beta_post+1))

sd_true = sqrt(var_true)

size = 10000
rtheta = rbeta(size,alpha_post,beta_post)

```

The code block above defined the parameters from Bernoulli model and beta prior. Due to beta prior is a conjugate prior, we can calculate the mean and variance by formula:

$$\mathbb{E}(\theta) = \frac{\alpha}{\alpha + \beta}$$

$$\mathbb{V}(\theta) = \frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)}$$

We use `rbeta()` to simulate a group of data from true posterior. According to Figure 1, as the number of sample size increases, the estimation of posterior mean converges to the true posterior mean, approximately  $\mathbb{E}(X) = 0.6667$ . Figure 2 also displayed the same tendency, where the standard deviation is very close to  $\sigma = 0.0943$

```

set.seed(42)
running_mean = 0
for (i in 1:size){
  running_mean[i] = mean(rtheta[1:i])
}

plot(running_mean, type='l',col="blue",lwd=2,xlab = "Numbers of Sample Size",
      ylab="Estimate of Posterior Mean")
abline(h=mean_true,col="red",lwd=2,lty=2)
legend("topright",legend=c("Monte-Carlo Estimate","True Posterior Mean"),col=c("blue","red"))

```

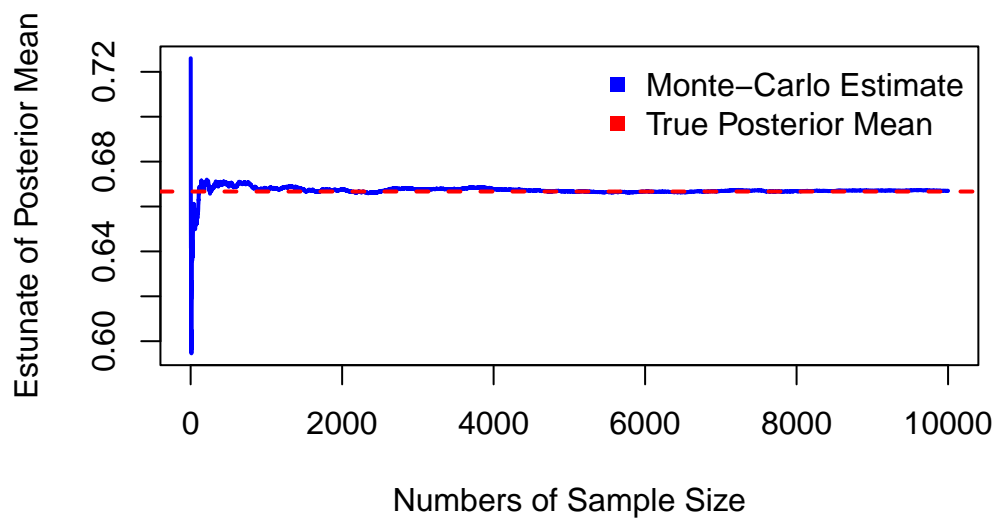


Figure 1: Convergence of True Posterior Mean by Monte-Carlo method

```
running_sd = 0
for (i in 1:size){
  running_sd[i] = sd(rtheta[1:i])
}

plot(running_sd, type='l',col="blue",lwd=2,xlab = "Numbers of Sample Size",
     ylab="Estimate of Posterior Standard Deviation")
abline(h=sd_true,col="red",lwd=2,lty=2)
legend("topright",legend=c("Monte-Carlo Estimate","True Posterior Standard Deviation"),col=c
```




Figure 2: Convergence of True Posterior Standard Deviation by Monte-Carlo method

### Problem 1b)

In this task, we calculate the posterior probability  $Pr(\theta < 0.5|y)$  by simulation, and compare the exact value by `pbeta()`. The result shows that the simulated answer is 0.0478, and the exact value is 0.0466, simulated value is pretty close to exact one.

```
nDraws = 10000
set.seed(42)

prob_sim = mean(rbeta(nDraws,alpha_post,beta_post)<=0.5)
prob_true = pbeta(0.5,alpha_post,beta_post)

# prob_sim = 0.0478
# prob_true 0.04656
```

### Problem 1c)

In this task, we simulate the posterior distribution of the log-odds  $\phi = \log(\frac{\theta}{1-\theta})$ . the method `qlogis()` can compute the data with log-odds transformation, which is equivalent to

```
transformed = log(theta/(1-theta)).
```

```
set.seed(42)
theta_original = rbeta(size,alpha_post,beta_post)
theta_trans = qlogis(theta_original)

hist(theta_trans,xlab="log-odds theta",main="",ylab="Density")
```




Figure 3: Histogram of log-odds theta

## Problem 2

In this task, we explore dataset `ericsson` on daily percentage returns on Ericsson stock. Figure 4 shows the distribution of standardized daily returns. Most values are centered around zero, but the distribution exhibits heavy tails, with occasional extreme negative returns.

```
load("ericsson.RData")
x = (returns - mean(returns))/sd(returns)
hist(x, 30, freq = FALSE, xlab = "daily returns (standardized)",
     ylab = "density",main = "")
```




Figure 4: Histogram of standardized daily returns

### Problem 2a)

We computed the log-likelihood function over a series of candidate degrees of freedom  $\nu$  and plotted the curve. In Figure 5, we can notice that the log-likelihood curve reaches its maximum value around 7, which implies that the maximum likelihood estimate of the degree of freedom is  $\hat{\nu} \approx 7$ .

```
nu_potential = seq(0.5,60,by=0.1)
loglike = numeric(length(nu_potential))

for (i in seq_along(nu_potential)){
  loglike[i] = sum(dt(x,df=nu_potential[i],log=TRUE))
}

plot(nu_potential,loglike,type='l',xlab="potential nu values",
     ylab="log-likelihood",lwd=2)

nu_mle = nu_potential[which.max(loglike)]
# nu_mle = 7
abline(v=nu_mle,col="red",lty=2,lwd=2)
```




Figure 5: Curve of log-likelihood in potential degrees of freedom  $\nu$

### Problem 2b)

We plot the likelihood  $L(x_1, \dots, x_n | \nu) = \prod_i p(x_i | \nu)$  and plotted the curve. in Figure 6, the red line and blue line represent  $\nu = 1$  and  $\nu = 10$  respectively. The likelihood peaks around  $\nu \approx 7$ . Clearly,  $L(1)$  (also called Cauchy Distribution) is obviously smaller than  $L(10)$ , showing that the data are heavy-tailed, but not as extreme as a Cauchy distribution.

```
nu_potential = seq(0.5,60,by=0.1)
likelihood = numeric(length(nu_potential))
for (i in seq_along(likelihood)){
  likelihood[i] = prod(dt(x,df=nu_potential[i],log=FALSE))
}
plot(nu_potential,likelihood,type='l',xlab="Potential nu values",
     ylab="Likelihood",lwd=2)
abline(v=1,col="red",lty=2,lwd=2)
abline(v=10,col="blue",lty=2,lwd=2)
legend("topright",c("nu = 1","nu = 10"),col = c("red","blue"),lty=2,lwd=2)
```




Figure 6: Curve of likelihood with respect to potential degrees of freedom  $\nu$

### Problem 2c)

In this step, we plot the logarithm of the posterior distribution for  $\nu$ , using

$$\log p(\nu \mid x_1, \dots, x_n) \propto \log p(x_1, \dots, x_n \mid \nu) + \log p(\nu),$$

we evaluate the log-likelihood over a series of candidate values of  $\nu$  and add the log prior. Note that the prior is  $\nu \sim \text{Exponential}(0.25)$  with the rate parameterization. The resulting curve (Figure 7) shows that the log-posterior peaks around  $\nu \approx 7$ .

```
nu_potential = seq(0.5,60,by=0.1)
log_post= numeric(length(nu_potential))
for (i in seq_along(nu_potential)){
  nu = nu_potential[i]
  logprior = dexp(nu,rate=0.25,log=TRUE)
  loglike = sum(dt(x,df=nu,log=TRUE))
  log_post[i] = loglike +logprior
}
plot(nu_potential,log_post,type='l',xlab='nu values',
     ylab = 'log-posterior',lwd=2)
```



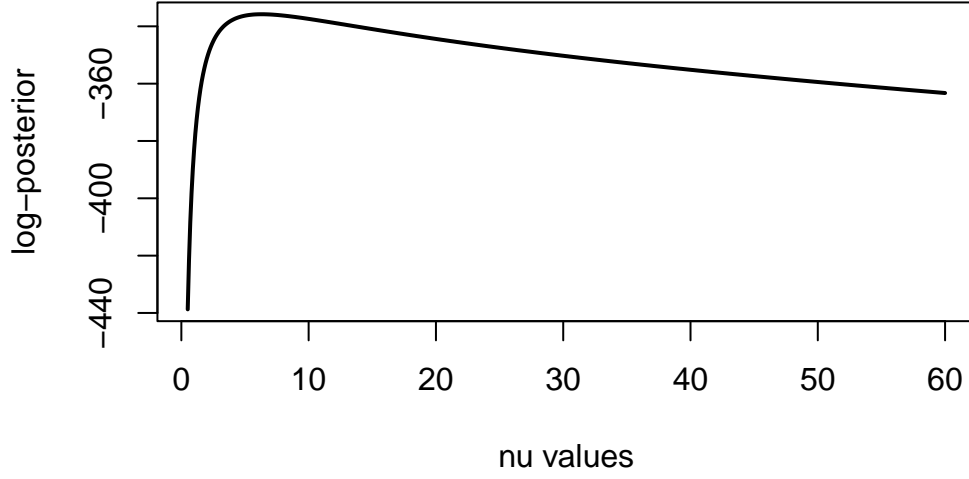


Figure 7: Curve of logarithm of the posterior distribution for nu

### Problem 2d)

In order to plot the posterior of distribution of  $nu$ , we firstly transform the log-posterior to unnormalized posterior as

$$p(\nu | x_1, \dots, x_n) \propto p(x_1, \dots, x_n | \nu)p(\nu)$$

Then, we need to normalize the posterior as a true probability density function:

$$p(\nu | x_1, \dots, x_n) = \frac{p(x_1, \dots, x_n | \nu)p(\nu)}{\int_0^\infty p(x_1, \dots, x_n | \nu)p(\nu)d\nu}$$

we can use Riemann approximation to calculate the integral:

$$p(\nu | x_1, \dots, x_n) \approx \frac{p(x_1, \dots, x_n | \nu_i)p(\nu_i)}{\sum_j p(x_1, \dots, x_n | \nu_j)p(\nu_j)\Delta\nu}$$

where  $\Delta\nu$  means the length of each step.

Figure 8 The blue line and orange line represents the posterior and prior respectively.

```

unnormalized_posterior = exp(log_post)
distance = 0.1
posterior = unnormalized_posterior/(sum(unnormalized_posterior)*distance)
plot(nu_potential,posterior, type="l",
     col = 'blue',xlab="nu",ylab="density",lwd=2)
lines(nu_potential,dexp(nu_potential,rate=0.25),col="orange",lwd=2)
legend("topright",c("posterior","prior"),col=c("blue","orange"),lty=1,lwd=4)

```




Figure 8: Posterior and Prior Distributions of the Degrees of Freedom  $\nu$

### Problem 2e)

The definition of Posterior mean is

$$\mathbb{E} \nu | x_1, \dots, x_n = \int_0^\infty \nu p(\nu | x_1, \dots, x_n) d\nu$$

where  $p(\nu | x_1, \dots, x_n)$  is normalized posterior distribution. We use Riemann sum, same as Problem 2d, to approximate the integral:

$$\mathbb{E}(\nu | x_1, \dots, x_n) \approx \sum_{i=1}^m \nu_i p(\nu_i | x_1, \dots, x_n) \Delta \nu$$

and the posterior mean of  $\nu = 7.0807$

```
post_mean = sum(nu_potential*posterior) * 0.1
```

```
library(mvtnorm)      # package with multivariate normal density
library(latex2exp)    # latex maths in plots
```

Warning: package 'latex2exp' was built under R version 4.4.3

## Problem 3

### Problem 3a)

The Gamma prior is conjugate to poisson model. We choose rate parameterization for gamma distribution. The posterior distribution is  $p(\theta|x_1, \dots, x_n) \sim \text{Gamma}(\alpha + n\bar{x}, \beta + n)$  where  $\alpha = 7$ ,  $\beta = 2$  from prior information.

```
## likelihood sample
sample_x = c(3, 5, 4, 3, 6, 8, 6, 1, 14, 3)
sample_x_size = length(sample_x)
sample_x_mean = mean(sample_x)
```

```
## parameters for prior
```

```
alpha = 7
beta = 2
```


```
# posterior simulation
```

```
n_draw = 10000
```

```
theta_sample = rgamma(n_draw, alpha+sum(sample_x), beta+sample_x_size)
```

```
hist(theta_sample, breaks = 50, probability = TRUE,
     col = "lightgray", border = "white",
     main = "Gamma Posterior Simulation",
     xlab = expression(theta))
```

## Gamma Posterior Simulation



```
## get all draw over 8
theta_sample_over_8 = theta_sample > 8
## calculate prob of theta>8 by using event_size/sample_size
sum(theta_sample_over_8)/n_draw
```

```
[1] 0
```

```
# use gamma cdf to get prob of over 8
pgamma(8, alpha+sum(sample_x), beta+sample_x_size, lower.tail = FALSE)
```

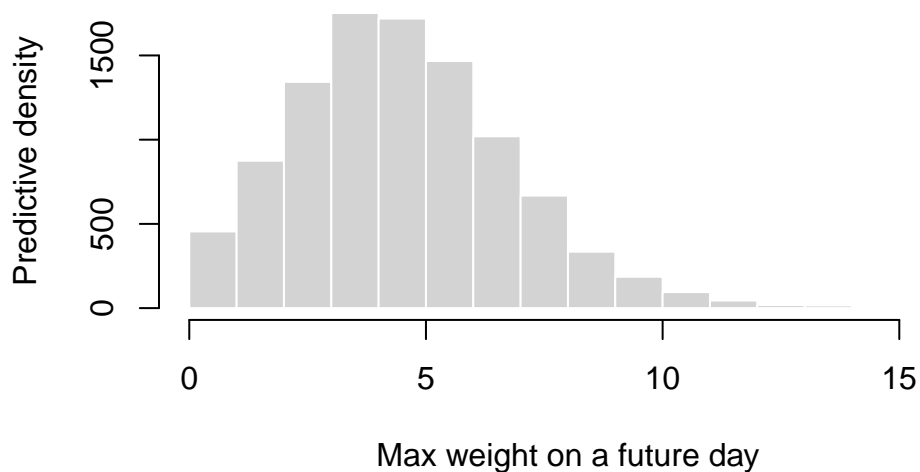
```
[1] 3.291869e-05
```

```
# check gamma pdf
# curve(dgamma(x, alpha+sum(sample_x), beta+sample_x_size), from=0, to=9,
# ylab="gamma pdf")
```

### Problem 3b)

```
# predictive simulation
predDraws = rnbino(n_draw, alpha+sum(sample_x),
                  (beta+sample_x_size)/(beta+sample_x_size+1))
hist(predDraws,
     xlab = "Max weight on a future day", ylab = "Predictive density",
     main = "Predictive density max weight - single day",
     col = "lightgray", border = "white")
```

### Predictive density max weight – single day



```
## get all draw over 8
temp = predDraws >= 8
## calculate prob of theta>=8 by using event_size/sample_size
mean(temp)
```

```
[1] 0.1369
```

```
# use gamma cdf to get prob of over 8
# negbinomial is discrete, cdf <= x, so using > 7
pnbinom(7, alpha+sum(sample_x), (beta+sample_x_size)/
        (beta+sample_x_size+1), lower.tail = FALSE)
```

```
[1] 0.141594
```

### Problem 3c)

The utility function is a function of random variable  $X_{11}$ . And  $a_{11}$  is treated as a constant.

$$U = \begin{cases} 10a_{11} & \text{if } X_{11} \geq a_{11}, \\ 10X_{11} - 7(a_{11} - X_{11}) & \text{if } X_{11} < a_{11} \end{cases} \quad (1)$$


With simplification, it should be

$$U = \begin{cases} 10a_{11} & \text{if } X_{11} \geq a_{11} \\ 17X_{11} - 7a_{11} & \text{if } X_{11} < a_{11} \end{cases} \quad (2)$$

The expected value of utility function is  $E(U) = 10a_{11} * Pr(X_{11} \geq a_{11}|a_{11}, x_1, \dots, x_{10}) + E(17X_{11} - 7a_{11}|X_{11} < a_{11}) * Pr(X_{11} < a_{11}|a_{11}, x_1, \dots, x_{10})$ . It is required to find the maximal value of expected utility. The uncertainty comes from demand  $X_{11}$  and storage  $a_{11}$ . The predictive distribution of  $X_{11}$  is from simulation. And only using  $a_{11}$  as variable and get the maximizer in expected utility.

```
aGrid = seq(0, 15, length = 1000)
EL = rep(length(aGrid))
for (i in 1:length(aGrid)){
  a = aGrid[i]
  p = mean(predDraws >= a)

  EL[i] = 10*a*p + (17*mean(predDraws[predDraws<a])-7*a)*(1-p)
}
plot(aGrid, EL, xlab = "storage, a", ylab = "Expected Utility", type = "l",
     lwd = 3, col = "lightgray")
abline(v = aGrid[which.max(EL)], lty = "dotted")
```



```
maximizer = aGrid[which.max(EL)]
maximizer
```

[1] 5

## Problem 4

Given information:

$$temp = \beta_0 + \beta_1 * time + \beta_2 * time^2 + \epsilon, \epsilon \sim N(0, \sigma^2)$$

```
#install.packages("remotes") # Uncomment this the first time
library(remotes)
#install_github("StatisticsSU/SUdatasets") # Uncomment this the first time
library(SUdatasets)
library(mvtnorm)
library(ggplot2)
head(tempLinkoping)
```

```

      time  temp
1 0.002732240  0.1
2 0.005464481 -4.5
3 0.008196721 -6.3
4 0.010928962 -9.6
5 0.013661202 -9.9
6 0.016393443 -17.1

```

```
summary(tempLinkoping)
```

```

      time          temp
Min.   :0.002732  Min.   : -17.100
1st Qu.:0.252049  1st Qu.:  1.925
Median :0.501366  Median :  6.900
Mean   :0.501366  Mean    :  7.524
3rd Qu.:0.750683  3rd Qu.: 14.575
Max.   :1.000000  Max.    : 23.100

```

```
cat("The dataset contains", length(tempLinkoping$time), "observations.")
```

The dataset contains 366 observations.

#### Problem 4a) Determine a suitable prior distribution

Given prior information:

$$\beta | \sigma^2 \sim N(\mu_0, \sigma^2 \Omega_0^{-1})$$

$$\sigma^2 \sim \text{inv-}\chi^2(\nu_0, \sigma_0^2)$$

The following figure shows the regression curves simulated from a beta prior with  $\mu_0(10, 100, -100)^T$ . From Figure 9, most of the temperatures variate above 0 degree which is higher than my belief of Linkoping's temperatures.

```

#beta prior ~ normal
mu0 = c(10,100,-100)
sigma_matrix0 = 0.01*diag(3)

#sigma_sq prior ~ inv-x^2

```



```

nu0 = 3
sigma0_sq = 1

#store simulated betas, sigma^2 and temps
m = 200
time_grid = tempLinkoping$time
temp_vals = matrix(NA, nrow=m, ncol=length(time_grid))

# Simulator for the scaled inverse Chi-square distribution
rScaledInvChi2 <- function(n, v_0, sigma2_0){
  return((v_0*sigma2_0)/rchisq(n, df = v_0))
}

#simulated sigma_prior and beta_prior
set.seed(42)

sigma_sq_draws = rScaledInvChi2(m, v_0=nu0, sigma2_0 = sigma0_sq)

for (i in 1:m){
  sigma_sq_i = sigma_sq_draws[i]
  betas_i = rmvnorm(1, mean=mu0, sigma=(sigma_sq_i*solve(sigma_matrix0)))
  temp_vals[i, ] = betas_i[1] + betas_i[2]*time_grid + betas_i[3]*time_grid^2
}

plot(
  NA, NA,
  xlim = c(0, 1),
  ylim = range(temp_vals),
  xlab = "Normalized Time",
  ylab = "Temperatures",
  main = "Simulated Regression Curves from Prior"
)

# Now overlay all the curves
for (i in 1:nrow(temp_vals)) {
  lines(time_grid, temp_vals[i, ], col = rgb(0, 0, 0, alpha = 0.3))
}

```




Figure 9: Simulated Regression Curves from Prior

The original prior information is much higher than my belief. Figure 10 is the regression curves plot with updated prior information  $\mu_0(5, 70, -70)^T$ , it clearly shows that half of the simulated temperature curves variate below 0 degree which satisfies my beliefs of temperature variation in Linköping.

```
#beta prior ~ normal updated
mu0 = c(5,70,-70)
Omega_matrix0 = 0.01*diag(3)

#sigma_sq prior ~ inv-x^2
nu0 = 3
sigma0_sq = 1

#store simulated betas, sigma^2 and temps
m = 200
time_grid = seq(0, 1, length.out=366)
temp_vals = matrix(NA, nrow=m, ncol=length(time_grid))

# Simulator for the scaled inverse Chi-square distribution
rScaledInvChi2 <- function(n, v_0, sigma2_0){
  return((v_0*sigma2_0)/rchisq(n, df = v_0))
}
```

```

}

#simulated sigma_prior and beta_prior
set.seed(42)

sigma_sq_draws = rScaledInvChi2(m, v_0=nu0, sigma2_0 = sigma0_sq)

for (i in 1:m){
  sigma_sq_i = sigma_sq_draws[i]
  betas_i = rmvnorm(1, mean=mu0, sigma=(sigma_sq_i*solve(Omega_matrix0)))
  temp_vals[i, ] = betas_i[1] + betas_i[2]*time_grid + betas_i[3]*time_grid^2
}

plot(
  NA, NA,
  xlim = c(0, 1),
  ylim = range(temp_vals),
  xlab = "Normalized Time",
  ylab = "Temperatures",
  main = "Simulated Regression Curves from Updated Prior"
)

# Now overlay all the curves
for (i in 1:nrow(temp_vals)) {
  lines(time_grid, temp_vals[i, ], col = rgb(0, 0, 0, alpha = 0.3))
}

```

## Simulated Regression Curves from Updated Prior




Figure 10: Simulated Regression Curves from Updated Prior

### Problem 4b) Simulating from the posterior

From the given information, a Gaussian linear regression with a conjugate prior will have a posterior with the same distribution family. The following is the joint posterior information.

$$\beta \mid \sigma^2, y \sim N(\mu_n, \sigma^2 \Omega_n^{-1})$$

$$\sigma^2 \mid y \sim Inv - \chi^2(\nu_n, \sigma_n^2)$$

$$\Omega_n = X^T X + \Omega_0$$

$$\mu_n = \Omega_n^{-1}(X^T y + \Omega_0 \mu_0)$$

$$\nu_n = \nu_0 + n$$

$$\sigma_n^2 = (\nu_0 \sigma_0^2 + y^T y + \mu_0^T \Omega_0 \mu_0 - \mu_n^T \Omega_n \mu_n) / \nu_n$$

The parameters information from prior and likelihood/model are given, we can compute the posterior parameters using the following code.

```

#settings before simulation
time = tempLinkoping$time
X = cbind(1, time, time^2)
y = tempLinkoping$temp
n = length(y)

#beta prior ~ normal updated
mu0 = c(5,70,-70)
Omega0 = 0.01*diag(3)

#sigma_sq prior ~ Inv-x^2
nu0 = 3
sigma0_sq = 1

#posterior settings
Omega_n = t(X) %*% X + Omega0
mu_n = solve(Omega_n) %*% (t(X) %*% y + Omega0 %*% mu0)
nu_n = nu0 + n
sigma_n_sq = (nu0*sigma0_sq + t(y)%*%y + t(mu0)%*%Omega0%*%mu0 -
               t(mu_n)%*%Omega_n%*%mu_n) / nu_n

```

After computing the posterior parameters, we can simulate samples from the posterior of  $\sigma^2$ . Then, the simulated samples of  $\sigma^2$  can be plugged into  $\beta \mid \sigma^2$ 's posterior distribution, which is multivariate normal. Finally, we can draw samples from  $\beta$ 's posterior distribution and visualize the marginal posteriors of each parameter using histograms as shown in Figure 11.

```

m = 10000
sigma_sq_post_draws = rScaledInvChi2(m, nu_n, sigma_n_sq)
#store betas draws
beta_post_draws = matrix(NA, nrow=m, ncol=ncol(X))

for (i in 1:m){
  sigma_sq_post_i = sigma_sq_post_draws[i]
  beta_post_var = sigma_sq_post_i * solve(Omega_n)

  beta_post_draws[i, ] = rmvnorm(1, mean=as.vector(mu_n), sigma=beta_post_var)
}

par(mfrow = c(2, 2))
beta_names = expression(beta[0], beta[1], beta[2])
beta_mains =c("Intercept", "Time", "Time^2")

```

```

for (i in 1:length(beta_names)){
  hist(beta_post_draws[, i], main=beta_mains[i], xlab=beta_names[i],
       col="skyblue",
       freq=FALSE)
}

hist(sigma_sq_post_draws, main="Variance", xlab=expression(sigma^2),
     col="skyblue", freq=FALSE)

```

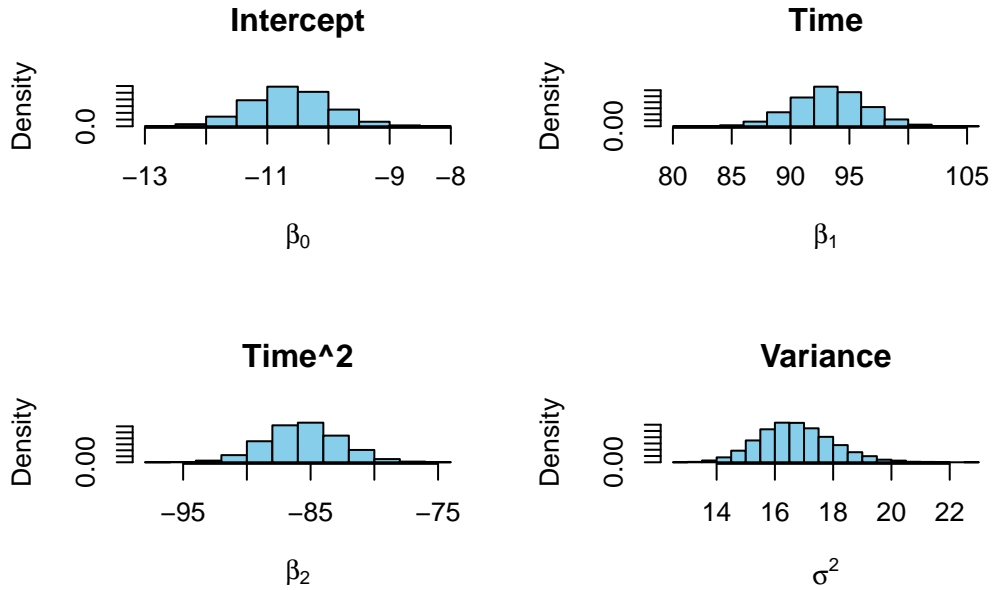


Figure 11: Marginal Distribution of Parameters

Since the marginal posterior distributions of coefficients  $\beta = (\beta_0, \beta_1, \beta_2)$  are known, the posterior distribution of the regression function can be computed by plugging the simulated  $\beta$  posterior draws into the regression model:

$$f(\text{time}) = \beta_0 + \beta_1 * \text{time} + \beta_2 * \text{time}^2$$

After the posterior distribution of the regression function is known, we can simply compute the posterior median and the equal-tail 95% confidence interval for the regression at each time point. Figure 12 contains the data point from Linköping dataset and a posterior median curve of the regression function with equal-tail 95% confidence interval.

```

# Predictive curves
time_grid = tempLinkoping$time
X_grid = cbind(1, time, time^2)
# X_grid: 366x3, beta_post_draws: 10000x3, transpose to 3x10000
f_post = X_grid %*% t(beta_post_draws)

f_median = apply(f_post, 1, median)
f_lower <- apply(f_post, 1, quantile, probs = 0.025)
f_upper <- apply(f_post, 1, quantile, probs = 0.975)

# Base R plot version of predictive regression curve

# Set up base plot
plot(time_grid, f_median, type = "l", lwd = 2, col = "red",
      ylim = range(c(f_lower, f_upper, y)),
      xlab = "Normalized Time", ylab = "Temperature",
      main = "Posterior Median of Regression Curve with 95% C.I")

# Add credible interval (ribbon)
polygon(c(time_grid, rev(time_grid)),
        c(f_lower, rev(f_upper)),
        col = rgb(70/255, 130/255, 180/255, 0.4), border = NA)

# Add observed data points
points(time, y, pch = 20, col = "gray40")

# Optionally add median line again (drawn over ribbon)
lines(time_grid, f_median, col = "red", lwd = 2)

```

### Posterior Median of Regression Curve with 95% C.I




Figure 12: Posterior Median of Regression Curve with 95% C.I

From the posterior median of the regression curve, we can clearly confirm that the 95% confidence interval band does not contain all the observed data point, because the confidence interval band only reflect the uncertainty of the posterior median regression function without the noise term  $\epsilon \sim N(0, \sigma^2)$ . Therefore, it is not necessary for the confidence interval band of the regression function to include every data points.

#### Problem 4c) Locating the day with the highest expected temperature

Given the highest expected temperature at each time point  $x_{max}$ :

$$x_{max} = -\frac{\beta_1}{2\beta_2}$$

Since the posterior distribution of  $\beta$  is known from previous problem, we can use the samples from the distribution to compute the values of  $x_{max}$  at each time point, and the values can be visualized through a histogram as shown in Figure 13.

```
beta1_post_samples = beta_post_draws[, 2]
beta2_post_samples = beta_post_draws[, 3]
```



```

x_max = -beta1_post_samples / (2*beta2_post_samples)

par(mfrow=c(1, 2))

hist(x_max, breaks=50, freq = FALSE, col="lightblue",
     main="Posterior of x_max",
     xlab="Time With Highest Expected Temperature")

#unnormalized time to regular day
hist(x_max*366, breaks=50, freq = FALSE, col="steelblue",
     main="Posterior of x_max",
     xlab="Day With Highest Expected Temperature")

```




Figure 13: Posterior of  $x_{\max}$